



PRIFYSGOL  
**BANGOR**  
UNIVERSITY

|                           |                       |
|---------------------------|-----------------------|
| <b>Student Name</b>       | <b>Gelareh Kabiri</b> |
| <b>Module Supervisor</b>  | <b>Heather He</b>     |
| <b>Module title</b>       | <b>Python coding</b>  |
| <b>Date of Submission</b> | <b>07/01/2024</b>     |

---

## Contents

|   |    |
|---|----|
| Introduction.....                                   | 3  |
| Task 1 .....  | 4  |
| Importing Data .....                                | 4  |
| Data Manipulation .....                             | 4  |
| Data Sampling.....                                  | 5  |
| Task 2 .....  | 6  |
| Job backgrounds and Education Levels Analysis ..... | 7  |
| Contact Communication Type Analysis.....            | 10 |
| Age Group Analysis .....                            | 12 |
| Relationship between Education and Responses .....  | 13 |
| Impact of Campaign Duration on Responses.....       | 14 |
| References.....                                     | 16 |

## Introduction

The Bank Marketing Dataset is a comprehensive resource that sheds light on the complex interaction between marketing and banking. It may be accessed through the UCI Machine Learning Repository at <https://archive.ics.uci.edu/dataset/222/bank+marketing>. This carefully selected dataset provides a thorough examination of consumer interactions in the banking industry, making it an invaluable resource for scholars, analysts, and business professionals. With a wide range of characteristics, the dataset offers insights into many facets of consumer behaviors and how they react to financial institutions' marketing campaigns. It contains information about job profiles, educational backgrounds, and demographics in addition to specifics regarding the communication channels that are employed in marketing efforts. Furthermore, economic indicators that help to provide a comprehensive picture of the economic environment in which these marketing initiatives take place include employment variation rates, consumer pricing indices, and current interest rates. The Bank Marketing Dataset is primarily intended for use in predictive modelling and analytics, but it also makes it easier to investigate patterns and trends that affect consumers' choices to sign up for term deposits. This data is crucial for developing strategic marketing plans, improving client interactions, and allocating resources as efficiently as possible in the banking industry. The UCI Machine Learning Repository hosts this publicly available dataset, which is a useful tool for data scientists, academic researchers, and business professionals. Its legitimacy and applicability for a range of uses, such as machine learning, statistical analysis, and exploratory data analysis, are highlighted by its availability.

# Task 1

## Importing Data

I imported the data for Task 1 and conducted the data analysis and manipulation procedures for the "Bank Marketing" dataset. The aim was to clean and analyze the data, then concentrate on taking a random 20% selection and analyzing it in further detail.

```
In [1]: 1 #Importing Liabraries
        2 import pandas as pd
        3 import numpy as nd

In [2]: 1 print("Date: 2024-12-26")
        2 print("Student ID: 500681622")
        3 print("Purpose: Python Assignment")
```

Date: 2024-12-26  
Student ID: 500681622  
Purpose: Python Assignment

## TASK\_1

```
In [3]: 1 #Importing dataset
        2 data= pd.read_csv("D:/University/Semester1/Coding/Python project/bank-additional-full.csv")
        3 print(data)

      age; "job"; "marital"; "education"; "default"; "housing"; "loan"; "contact"; "month"; "day_of_week"; "duration"; "campaign"; "pday
s"; "previous"; "poutcome"; "emp.var.rate"; "cons.price.idx"; "cons.conf.idx"; "euribor3m"; "nr.employed"; "y"
0      56; "housemaid"; "married"; "basic.4y"; "no"; "no"; ...
1      57; "services"; "married"; "high.school"; "unknown...
2      37; "services"; "married"; "high.school"; "no"; "ye...
3      40; "admin."; "married"; "basic.6y"; "no"; "no"; "no...
4      56; "services"; "married"; "high.school"; "no"; "no...
...
41183  73; "retired"; "married"; "professional.course"; "...
41184  46; "blue-collar"; "married"; "professional.cours...
41185  56; "retired"; "married"; "university.degree"; "no...
41186  44; "technician"; "married"; "professional.course...
41187  74; "retired"; "married"; "professional.course"; "...

[41188 rows x 1 columns]
```

## Data Manipulation

It was noticed that the dataset required more processing after importing it since it contained values that were semicolon separated. It was fixed by changing the read\_csv() function to use the right delimiter, which produced a cleaner dataset. Delimiters are used in literature to organize the data set for processing (Miller, 2018). To understand the structure of the dataset and spot any potential

problems, including errors or missing information, I also looked over it. To make sure a thorough comprehension of the data is essential.

```
In [4]: 1 print("data preview:")
        2 print(data.head())
```

data preview:

```
age;job;marital;education;default;housing;loan;contact;month;day_of_week;duration;campaign;pdays;previous;poutcome;emp.var.rate;cons.price.idx;cons.conf.idx;euribor3m;nr.employed;y
0 56;housemaid;married;basic.4y;no;no;no;telephone;may;mon...
1 57;services;married;high.school;unknown;no;no;telephone;
2 37;services;married;high.school;no;yes;no;telephone;
3 40;admin.;married;basic.6y;no;no;no;telephone;
4 56;services;married;high.school;no;no;yes;telephone;
```

```
In [5]: 1 # Data_Manipulation
        2 data= pd.read_csv("D:/University/Semester1/Coding/Python project/bank-additional-full.csv", delimiter=";", quotechar='')
        3 print(data.head())
```

|   | age | job       | marital | education   | default | housing | loan | contact   | \ |
|---|-----|-----------|---------|-------------|---------|---------|------|-----------|---|
| 0 | 56  | housemaid | married | basic.4y    | no      | no      | no   | telephone |   |
| 1 | 57  | services  | married | high.school | unknown | no      | no   | telephone |   |
| 2 | 37  | services  | married | high.school | no      | yes     | no   | telephone |   |
| 3 | 40  | admin.    | married | basic.6y    | no      | no      | no   | telephone |   |
| 4 | 56  | services  | married | high.school | no      | no      | yes  | telephone |   |

  

|   | month | day_of_week | ... | campaign | pdays | previous | poutcome    | emp.var.rate | \ |
|---|-------|-------------|-----|----------|-------|----------|-------------|--------------|---|
| 0 | may   | mon         | ... | 1        | 999   | 0        | nonexistent | 1.1          |   |
| 1 | may   | mon         | ... | 1        | 999   | 0        | nonexistent | 1.1          |   |
| 2 | may   | mon         | ... | 1        | 999   | 0        | nonexistent | 1.1          |   |
| 3 | may   | mon         | ... | 1        | 999   | 0        | nonexistent | 1.1          |   |
| 4 | may   | mon         | ... | 1        | 999   | 0        | nonexistent | 1.1          |   |

  

|   | cons.price.idx | cons.conf.idx | euribor3m | nr.employed | y  |
|---|----------------|---------------|-----------|-------------|----|
| 0 | 93.994         | -36.4         | 4.857     | 5191.0      | no |
| 1 | 93.994         | -36.4         | 4.857     | 5191.0      | no |
| 2 | 93.994         | -36.4         | 4.857     | 5191.0      | no |
| 3 | 93.994         | -36.4         | 4.857     | 5191.0      | no |
| 4 | 93.994         | -36.4         | 4.857     | 5191.0      | no |

[5 rows x 21 columns]

## Data Sampling

A random sample of 20% of the cleaned dataset was taken to make further studies easier. To guarantee that the sampling procedure could be repeated, a random seed was chosen. To extract a random subset, the `random_state` parameter was set to 42 for reproducibility using the `sample()` method.

```
In [6]: 1 ## Data_sampling
2 random_seed = 42
3 sampled_data = data.sample(frac=0.2, random_state=random_seed)
4 print("Sampled data preview:")
5 print(sampled_data.head())
```

Sampled data preview:

|       | age | job         | marital | education   | default | housing | loan | \ |
|-------|-----|-------------|---------|-------------|---------|---------|------|---|
| 32884 | 57  | technician  | married | high.school | no      | no      | yes  |   |
| 3169  | 55  | unknown     | married | unknown     | unknown | yes     | no   |   |
| 32206 | 33  | blue-collar | married | basic.9y    | no      | no      | no   |   |
| 9403  | 36  | admin.      | married | high.school | no      | no      | no   |   |
| 14020 | 27  | housemaid   | married | high.school | no      | yes     | no   |   |

  

|       | contact   | month | day_of_week | ... | campaign | pdays | previous | \ |
|-------|-----------|-------|-------------|-----|----------|-------|----------|---|
| 32884 | cellular  | may   | mon         | ... | 1        | 999   | 1        |   |
| 3169  | telephone | may   | thu         | ... | 2        | 999   | 0        |   |
| 32206 | cellular  | may   | fri         | ... | 1        | 999   | 1        |   |
| 9403  | telephone | jun   | fri         | ... | 4        | 999   | 0        |   |
| 14020 | cellular  | jul   | fri         | ... | 2        | 999   | 0        |   |

  

|       | poutcome    | emp.var.rate | cons.price.idx | cons.conf.idx | euribor3m | \ |
|-------|-------------|--------------|----------------|---------------|-----------|---|
| 32884 | failure     | -1.8         | 92.893         | -46.2         | 1.299     |   |
| 3169  | nonexistent | 1.1          | 93.994         | -36.4         | 4.860     |   |
| 32206 | failure     | -1.8         | 92.893         | -46.2         | 1.313     |   |
| 9403  | nonexistent | 1.4          | 94.465         | -41.8         | 4.967     |   |
| 14020 | nonexistent | 1.4          | 93.918         | -42.7         | 4.963     |   |

  

|       | nr.employed | y  |
|-------|-------------|----|
| 32884 | 5099.1      | no |
| 3169  | 5191.0      | no |
| 32206 | 5099.1      | no |
| 9403  | 5228.1      | no |
| 14020 | 5228.1      | no |

[5 rows x 21 columns]

## Task 2

In Task 2, the "Bank Marketing" dataset was explored and analysed to use the proper data analytics approaches to answer queries. Using a classification system to address important questions was the focus of the investigation.

## Job backgrounds and Education Levels Analysis

### Job Backgrounds and Education Levels Analysis

```
In [7]: 1 # Assuming 'job' and 'education' are relevant features for analysis
2 X1 = sampled_data[['job', 'education']]
3 y1 = sampled_data['y']
4
5 # Splitting the data into training and testing sets
6 from sklearn.model_selection import train_test_split
7
8 X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=42)
9
10 # One-hot encoding for categorical variables
11 from sklearn.preprocessing import OneHotEncoder
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14
15 # Define the transformer for one-hot encoding
16 preprocessor1 = ColumnTransformer(
17     transformers=[
18         ('cat', OneHotEncoder(), ['job', 'education'])
19     ],
20     remainder='passthrough'
21 )
22
23 # Assuming Logistic Regression as the chosen classification algorithm
24 from sklearn.linear_model import LogisticRegression
25 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
26
27 # Create a pipeline with one-hot encoding and Logistic regression
28 model1 = Pipeline(steps=[
29     ('preprocessor', preprocessor1),
30     ('classifier', LogisticRegression(random_state=42))
31 ])
32
33 # Training the model
34 model1.fit(X1_train, y1_train)
35
36 # Predicting responses on the test set
37 predictions1 = model1.predict(X1_test)
38
39 # Evaluating the model
40 accuracy1 = accuracy_score(y1_test, predictions1)
41 conf_matrix1 = confusion_matrix(y1_test, predictions1)
42 classification_report1 = classification_report(y1_test, predictions1)
43
44 # Displaying results
45 print(f"Accuracy: {accuracy1}")
46 print("\nConfusion Matrix:\n", conf_matrix1)
47 print("\nClassification Report:\n", classification_report1)
```



```

Accuracy: 0.8743932038834952

Confusion Matrix:
[[1441   0]
 [ 207   0]]

Classification Report:

```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| no           | 0.87      | 1.00   | 0.93     | 1441    |
| yes          | 0.00      | 0.00   | 0.00     | 207     |
| accuracy     |           |        | 0.87     | 1648    |
| macro avg    | 0.44      | 0.50   | 0.47     | 1648    |
| weighted avg | 0.76      | 0.87   | 0.82     | 1648    |

A machine learning pipeline for response prediction based on "job" and "education"-related features is built in the code that is provided. These are the logical actions performed: The `train_test_split` function from scikit-learn is used to divide the data into training and testing sets. Column Transformer is used to incorporate the one-hot encoded categorical variables ('job' and 'education') into the preprocessing pipeline. Scikit-learn's Pipeline is used to design a machine learning pipeline. The Pipeline tool from Scikit-learn is a comprehensive Python utility that simplifies the process of building and running machine learning processes. Müller and Guido (2017) describe this capability as allowing users to easily combine various machine learning models and data pretreatment procedures into a cohesive and effective pipeline. The 'Pipeline' class streamlines the code structure and improves repeatability by encapsulating different phases of data processing, including data transformation and model training. This methodology ensures a well-organized and systematic approach to machine learning tasks, while also encouraging cleaner and more understandable code and the consistent implementation of a series of operations on the data. A logistic regression classifier and a preprocessor manage one-hot encoding in this pipeline. Using the pipeline's fit approach, the logistic regression model is trained on the training set. The trained model is used to make predictions on the test set. Metrics from the classification report, accuracy, and confusion matrix are used to assess the model's performance. A machine learning model's accuracy, as an indicator of overall correctness, is determined by the ratio of correctly predicted instances to the total instances. The computation of accuracy can be accomplished using the 'accuracy\_score' function from the scikit-learn library in Python (Haghighi et al., 2018). Additionally, the performance of a classification method is visually summarized through a confusion matrix, which showcases the count of false positives, false negatives, true positives, and true negatives. The generation of this matrix in Python is facilitated by the 'confusion\_matrix' function in the scikit-learn library (Haghighi et al., 2018).



Part 1's analysis clarifies the substantial influence that customers' educational backgrounds and work histories have on their response rates. Based on the occupational histories and educational levels of the consumers, the study predicts their responses with a reasonable degree of accuracy (87.44% overall). This is achieved using a logistic regression model. When these elements are examined more closely, subtle patterns in their influence are revealed by the model's performance measures. The model demonstrates competence in anticipating negative answers ('no') in a range of educational and professional backgrounds. Its capacity to forecast affirmative answers ('yes') shows unpredictability, though. The 'yes' class's precision, recall, and F1-score metrics shed light on the model's capacity to correctly identify favourable outcomes depending on educational and occupational backgrounds. The results essentially point to a significant correlation between the response rates of the clients and their educational and professional backgrounds. The significance of job background and education level in affecting consumers' likely to reply positively is highlighted by the model's ability to predict positive responses across several categories. To gain a more thorough grasp of these dynamics, greater investigation and refining may provide deeper insights into the precise effects of educational attainment and employment histories on customers' response rates.

## Contact Communication Type Analysis

### Contact Communication Type Analysis

```
In [8]: 1 # Assuming 'contact' is a relevant feature for analysis
2 X2 = sampled_data[['contact']]
3 y2 = sampled_data['y']
4
5 # Splitting the data into training and testing sets
6 from sklearn.model_selection import train_test_split
7
8 X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.2, random_state=42)
9
10 # One-hot encoding for the 'contact' column
11 from sklearn.preprocessing import OneHotEncoder
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14
15 # Define the transformer for one-hot encoding
16 preprocessor2 = ColumnTransformer(
17     transformers=[
18         ('cat', OneHotEncoder(), ['contact'])
19     ],
20     remainder='passthrough'
21 )
22
23 # Assuming Logistic Regression as the chosen classification algorithm
24 from sklearn.linear_model import LogisticRegression
25 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
26
27 # Create a pipeline with one-hot encoding and logistic regression
28 model2 = Pipeline(steps=[
29     ('preprocessor', preprocessor2),
30     ('classifier', LogisticRegression(random_state=42))
31 ])
32
33 # Training the model
34 model2.fit(X2_train, y2_train)
35
36 # Predicting responses on the test set
37 predictions2 = model2.predict(X2_test)
38
39 # Evaluating the model
40 accuracy2 = accuracy_score(y2_test, predictions2)
41 conf_matrix2 = confusion_matrix(y2_test, predictions2)
42 classification_report2 = classification_report(y2_test, predictions2)
43
44 # Displaying results
45 print(f"Accuracy: {accuracy2}")
46 print("\nConfusion Matrix:\n", conf_matrix2)
47 print("\nClassification Report:\n", classification_report2)
```

Accuracy: 0.8743932038834952

Confusion Matrix:

```
[[1441   0]
 [ 207   0]]
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| no           | 0.87      | 1.00   | 0.93     | 1441    |
| yes          | 0.00      | 0.00   | 0.00     | 207     |
| accuracy     |           |        | 0.87     | 1648    |
| macro avg    | 0.44      | 0.50   | 0.47     | 1648    |
| weighted avg | 0.76      | 0.87   | 0.82     | 1648    |

The presented code involves multiple important phases: To begin with, the target variable 'y' was created, and the 'contact' column was chosen for analysis. The next steps were taken like previous question to build the model. This section's analysis shows a clear correlation between the selected contact communication kinds and consumers' response rates. The total accuracy of the logistic regression model used to evaluate this association is about 87.44%. This precision illustrates how well the model works to forecast consumers' reactions depending on the designated contact communication kinds. But a deeper look at the model's performance indicators reveals a significant drawback. High precision, recall, and F1-score metrics show that the model is quite good at predicting negative responses ('no'); nevertheless, it has difficulty correctly predicting positive responses ('yes'). The model may not be able to identify and predict positive outcomes based on contact communication types, as evidenced by the consistently low precision, recall, and F1-score for the 'yes' class throughout analyses. Although there is a clear correlation between the types of contact communications and consumers' response rates, the model's shortcomings in anticipating favourable replies call for additional research and possible improvement.

## Age Group Analysis

### Age Groups Investigation

```
In [9]: 1 # Assuming 'age' and 'contact' are relevant features for analysis
2 X3 = sampled_data[['age', 'contact']]
3 y3 = sampled_data['y']
4
5 # Splitting the data into training and testing sets
6 from sklearn.model_selection import train_test_split
7
8 X3_train, X3_test, y3_train, y3_test = train_test_split(X3, y3, test_size=0.2, random_state=42)
9
10 # One-hot encoding for the 'contact' column
11 from sklearn.preprocessing import OneHotEncoder
12 from sklearn.compose import ColumnTransformer
13 from sklearn.pipeline import Pipeline
14
15 # Define the transformer for one-hot encoding
16 preprocessor3 = ColumnTransformer(
17     transformers=[
18         ('cat', OneHotEncoder(), ['contact'])
19     ],
20     remainder='passthrough'
21 )
22
23 # Assuming Logistic Regression as the chosen classification algorithm
24 from sklearn.linear_model import LogisticRegression
25 from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
26
27 # Create a pipeline with one-hot encoding and Logistic regression
28 model3 = Pipeline(steps=[
29     ('preprocessor', preprocessor3),
30     ('classifier', LogisticRegression(random_state=42))
31 ])
32
33 # Training the model
34 model3.fit(X3_train, y3_train)
35
36 # Predicting responses on the test set
37 predictions3 = model3.predict(X3_test)
38
39 # Evaluating the model
40 accuracy3 = accuracy_score(y3_test, predictions3)
41 conf_matrix3 = confusion_matrix(y3_test, predictions3)
42 classification_report3 = classification_report(y3_test, predictions3)
43
44 # Displaying results
45 print(f"Accuracy: {accuracy3}")
46 print("\nConfusion Matrix:\n", conf_matrix3)
47 print("\nClassification Report:\n", classification_report3)
```

Accuracy: 0.8743932038834952

Confusion Matrix:

```
[[1441   0]
 [ 207   0]]
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| no           | 0.87      | 1.00   | 0.93     | 1441    |
| yes          | 0.00      | 0.00   | 0.00     | 207     |
| accuracy     |           |        | 0.87     | 1648    |
| macro avg    | 0.44      | 0.50   | 0.47     | 1648    |
| weighted avg | 0.76      | 0.87   | 0.82     | 1648    |

The 'age' and 'contact' features are the focus of the code when analysing client reaction. The following steps were taken just like the two previous sets of codes. An analysis of the Age Group Investigation using the logistic regression model indicates a consistent link between contact communication type and answers among different age groups. Regardless of age demographics, the general accuracy of roughly 87.44% is consistent. The consistent results observed in the confusion matrix and classification report correspond to the same patterns for the precision, recall, and F1-score metrics for both 'yes' and 'no' replies across various age categories. But as we can see from the precision, recall, and F1-score values of 0 for the 'yes' class, the model has trouble correctly predicting affirmative replies ('yes') across all age groups. Even while the contact communication type - replies association is consistent, the model's ability to detect good outcomes consistently lags. Even though the association holds true for all age groups, the model's inability to accurately predict positive responses indicates that more research and development are necessary to improve its accuracy in predicting good results in contact communication types for all age groups.

## Relationship between Education and Responses

```
In [10]: 1 #Relationship between education and responses
2 education_insight = sampled_data.groupby(['education', 'y']).size().unstack(fill_value=0)
3 print("\nInsight 1: Relationship between education and responses")
4 print(education_insight)
```

```
Insight 1: Relationship between education and responses
y
education
basic.4y      755   83
basic.6y      401   39
basic.9y     1144   93
high.school  1684  217
illiterate      4    1
professional.course  943  124
university.degree 2100  318
unknown       272   60
```

The given code segment examines the relationship between customer reactions and educational attainment. First, the 'education' and 'y' (customer replies) columns are used to group the dataset. The next action is to tally the instances of every distinct combination. The final data is then

reorganised into a more understandable pivot table style, with counts acting as the intersections, 'education' values acting as the index, and 'y' values acting as columns. When handling missing data, the `fill_value=0` argument is used to replace the missing values with zeros. The code concludes by printing the acquired insights to the console, providing a concise and organised summary of the relationship between various educational levels and client reactions. The analysis provides interesting new information on how responses to the marketing effort and educational level are related. Upon analysis of the data, we find that response patterns vary throughout various educational classifications. With 318 good responses against 2100 negative, customers with a university degree have the highest positive response rate. Though not as much as those with a university degree, those with professional courses also show a positive reaction trend. On the other hand, customers with a basic education background, such as basic.4y, fundamental 6y, and basic.9y tend to display more negative than positive answers. Despite having a very small sample size, the illiterate category receives a positive response. Still, this category has few data points, thus interpretation should be done with caution. Higher education levels are typically linked to more positive answers, while lower education levels show a more mixed pattern. It suggests that education level influences customers' responses. Given the importance of these developments, more research is necessary to fully comprehend the relationships between campaign outcomes and education.

## Impact of Campaign Duration on Responses

```

6 #Impact of campaign duration on responses
7 sampled_data['duration_category'] = pd.cut(sampled_data['duration'], bins=[0, 100, 200, 300, 400, 500, float('inf')], labels
8 duration_insight = sampled_data.groupby(['duration_category', 'y']).size().unstack(fill_value=0)
9 print("\nInsight 2: Impact of campaign duration on responses")
10 print(duration_insight)

```

```

Insight 2: Impact of campaign duration on responses
y
duration_category    no  yes
0-100                1995   11
101-200              2388  128
201-300              1277  155
301-400               696  116
401-500               366   88
501+                  581  437

```

The study in this code focuses on figuring out how the length of the campaign affects responses. Creating discrete time categories, classifying the data according to these categories and response

outcomes, and tabulating the results are the tasks involved. First, the dataset gains a new column called "duration\_category." Next, bins 0-100, 101-200, 201-300, 301-400, 401-500, and 501+ are created based on the 'duration' column. The `pd.cut` function is used to carry out this categorization. According to Unpingco (2021), the '`pd.cut`' method in Python is a potent tool in the pandas library that is used to bin numerical data into discrete intervals. This technique makes it easier to analyse distribution patterns and trends within various ranges by allowing users to split and classify a continuous variable into predetermined bins. It provides a versatile and effective technique to alter and modify data for statistical and exploratory purposes. It is especially helpful when working with datasets having numerical values that need to be categorised for different analyses. The dataset is then categorised according to the 'duration\_category' and the response target variable, 'y'. These groupings are made using the `groupby` method. Data can be grouped by particular columns in Python with the help of the pandas library's '`groupby`' function. It helps with operations like aggregation, transformation, and filtering by enabling the deployment of functions to each group independently. This technique is crucial for quickly examining patterns and trends within dataset subsets, offering insightful information on the distribution and properties of the data (Slatkin, 2019). Next, the `size` method is used to aggregate the data and find the number of occurrences for each group. The data is reshaped using the `unstack` approach to display the results in a way that is easier to understand. With this transformation, the aggregated data is pivoted, with 'y' representing the rows and 'duration\_category' the columns. Zeros are used to fill in any missing values (`fill_value=0`). The code then outputs the insights it has collected, paying particular attention to how campaign duration affects answers. It is possible to clearly see how various duration categories link to the associated response results thanks to the tabular style. There are only 11 positive responses out of 1995 cases in the '0-100' duration category, with the bulk of responses being negative. Positive replies rise dramatically to 128 as the period approaches the '101-200' range, while negative responses stay higher at 2388. The positive answers rise to 155 in the '201-300' length category as well, suggesting that longer campaign durations may have a beneficial impact. With 1277 occurrences, negative comments still outweigh positive ones even in this range. With differing degrees of positive and negative responses, the tendency is maintained in the ensuing length categories. The '501+' category is particularly noteworthy as it exhibits a significant rise in positive replies (437) relative to negative responses (581). Longer campaign durations may



be associated with more favourable replies, according to the results; however, the frequency of negative responses across all duration categories points to a complex link between campaign duration and customer involvement.

## References

archive.ics.uci.edu. (n.d.). *UCI Machine Learning Repository*. [online] Available at: <https://archive.ics.uci.edu/dataset/222/bank+marketing>.

Galli, S. (2020). *Python feature engineering cookbook: over 70 recipes for creating, engineering, and transforming features to build machine learning models*. Birmingham, UK: Packt Publishing.

Haghighi, S., Jasemi, M., Hessabi, S. and Zolanvari, A. (2018). PyCM: Multiclass confusion matrix library in Python. *Journal of Open-Source Software*, 3(25), p.729. doi: <https://doi.org/10.21105/joss.00729>.

Miller, C. (2018). *Hands-on data analysis with NumPy and pandas: implement Python packages from data manipulation to processing*. Birmingham: Packt Publishing Ltd.

Müller, A.C. and Guido, S. (2017). *Introduction to machine learning with Python: a guide for data scientists*. Beijing: O'reilly.

Raúl Garreta, Moncecchi, G., Hauck, T. and Hackeling, G. (2017). *Scikit-learn: machine learning simplified*. Birmingham, Uk: Packt Publishing.

Slatkin, B. (2019). *Effective Python*. Addison-Wesley Professional.

Unpingco, J. (2021). *Python Programming for Data Analysis*. Springer Nature.

## Appendix

### Codes

*#Importing liabraries*

```
import pandas as pd
```

```
import numpy as nd
```

In [2]:

```
print("Date: 2024-12-26")
```

```
print("Student ID: 500681622")
```

```
print("Purpose: Python Assignment")
```

Date: 2024-12-26

Student ID: 500681622

Purpose: Python Assignment

### TASK\_1

In [3]:

*#Importing dataset*

```
data= pd.read_csv("D:/University/Semester1/Coding/Python project/bank-additional-full.csv")
```

```
print(data)
```

```
age;"job";"marital";"education";"default";"housing";"loan";"contact";"month";"day_of_week";"d
uration";"campaign";"pdays";"previous";"poutcome";"emp.var.rate";"cons.price.idx";"cons.conf.
idx";"euribor3m";"nr.employed";"y"
```

```
0    56;"housemaid";"married";"basic.4y";"no";"no";...
```

```
1    57;"services";"married";"high.school";"unknown...
```

```
2    37;"services";"married";"high.school";"no";"ye...
```

```
3    40;"admin."; "married";"basic.6y";"no";"no";"no...
```

```
4    56;"services";"married";"high.school";"no";"no...
```

```
...                               ...
```

```
41183 73;"retired";"married";"professional.course";"...
```

```
41184 46;"blue-collar";"married";"professional.cours...
```

```
41185 56;"retired";"married";"university.degree";"no...
```

```
41186 44;"technician";"married";"professional.course...
```

```
41187 74;"retired";"married";"professional.course";"...
```

```
[41188 rows x 1 columns]
```

```
In [4]:
```

```
print("data preview:")
```

```
print(data.head())
```

data preview:

```
age;"job";"marital";"education";"default";"housing";"loan";"contact";"month";"day_of_week";"d
uration";"campaign";"pdays";"previous";"poutcome";"emp.var.rate";"cons.price.idx";"cons.conf.
idx";"euribor3m";"nr.employed";"y"
```

```
0 56;"housemaid";"married";"basic.4y";"no";"no";...
```

```
1 57;"services";"married";"high.school";"unknown...
```

```
2 37;"services";"married";"high.school";"no";"ye...
```

```
3 40;"admin."; "married";"basic.6y";"no";"no";"no...
```

```
4 56;"services";"married";"high.school";"no";"no...
```

In [5]:

```
# Data_Manipulation
```

```
data= pd.read_csv("D:/University/Semester1/Coding/Python project/bank-additional-full.csv",
delimiter=";", quotechar="")
```

```
print(data.head())
```

```
age    job  marital  education  default  housing  loan  contact \
0  56  housemaid  married  basic.4y    no     no   no  telephone
1  57  services  married  high.school  unknown  no     no  telephone
2  37  services  married  high.school    no    yes   no  telephone
```

3 40 admin. married basic.6y no no no telephone

4 56 services married high.school no no yes telephone

month day\_of\_week ... campaign pdays previous poutcome emp.var.rate \

0 may mon ... 1 999 0 nonexistent 1.1

1 may mon ... 1 999 0 nonexistent 1.1

2 may mon ... 1 999 0 nonexistent 1.1

3 may mon ... 1 999 0 nonexistent 1.1

4 may mon ... 1 999 0 nonexistent 1.1

cons.price.idx cons.conf.idx euribor3m nr.employed y

0 93.994 -36.4 4.857 5191.0 no

1 93.994 -36.4 4.857 5191.0 no

2 93.994 -36.4 4.857 5191.0 no

3 93.994 -36.4 4.857 5191.0 no

4 93.994 -36.4 4.857 5191.0 no

[5 rows x 21 columns]

In [6]:

```
## Data_sampling
```

```
random_seed = 42
```

```
sampled_data = data.sample(frac=0.2, random_state=random_seed)
```

```
print("Sampled data preview:")
```

```
print(sampled_data.head())
```

Sampled data preview:

|       | age | job         | marital | education   | default | housing | loan \ |
|-------|-----|-------------|---------|-------------|---------|---------|--------|
| 32884 | 57  | technician  | married | high.school | no      | no      | yes    |
| 3169  | 55  | unknown     | married | unknown     | unknown | yes     | no     |
| 32206 | 33  | blue-collar | married | basic.9y    | no      | no      | no     |
| 9403  | 36  | admin.      | married | high.school | no      | no      | no     |
| 14020 | 27  | housemaid   | married | high.school | no      | yes     | no     |

|       | contact   | month | day_of_week | ... | campaign | pdays | previous \ |
|-------|-----------|-------|-------------|-----|----------|-------|------------|
| 32884 | cellular  | may   | mon         | ... | 1        | 999   | 1          |
| 3169  | telephone | may   | thu         | ... | 2        | 999   | 0          |
| 32206 | cellular  | may   | fri         | ... | 1        | 999   | 1          |

9403 telephone jun fri ... 4 999 0

14020 cellular jul fri ... 2 999 0

poutcome emp.var.rate cons.price.idx cons.conf.idx euribor3m \

32884 failure -1.8 92.893 -46.2 1.299

3169 nonexistent 1.1 93.994 -36.4 4.860

32206 failure -1.8 92.893 -46.2 1.313

9403 nonexistent 1.4 94.465 -41.8 4.967

14020 nonexistent 1.4 93.918 -42.7 4.963

nr.employed y

32884 5099.1 no

3169 5191.0 no

32206 5099.1 no

9403 5228.1 no

14020 5228.1 no

[5 rows x 21 columns]



## TASK\_2

### Job Backgrounds and Education Levels Analysis

In [7]:

```
# Assuming 'job' and 'education' are relevant features for analysis
```

```
X1 = sampled_data[['job', 'education']]
```

```
y1 = sampled_data['y']
```

```
# Splitting the data into training and testing sets
```

```
from sklearn.model_selection import train_test_split
```

```
X1_train, X1_test, y1_train, y1_test = train_test_split(X1, y1, test_size=0.2, random_state=42)
```

```
# One-hot encoding for categorical variables
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.pipeline import Pipeline
```

```
# Define the transformer for one-hot encoding
```

```
preprocessor1 = ColumnTransformer(
```

```
    transformers=[
```

```
        ('cat', OneHotEncoder(), ['job', 'education'])
```

```
    ],
```

```
    remainder='passthrough'
```

```
)
```

```
# Assuming Logistic Regression as the chosen classification algorithm
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Create a pipeline with one-hot encoding and logistic regression
```

```
model1 = Pipeline(steps=[
```

```
    ('preprocessor', preprocessor1),
```

```
    ('classifier', LogisticRegression(random_state=42))
```

```
])
```

```
# Training the model
```

```
model1.fit(X1_train, y1_train)
```

```
# Predicting responses on the test set
```

```
predictions1 = model1.predict(X1_test)
```

```
# Evaluating the model
```

```
accuracy1 = accuracy_score(y1_test, predictions1)
```

```
conf_matrix1 = confusion_matrix(y1_test, predictions1)
```

```
classification_report1 = classification_report(y1_test, predictions1)
```

```
# Displaying results
```

```
print(f'Accuracy: {accuracy1}')
```

```
print("\nConfusion Matrix:\n", conf_matrix1)
```

```
print("\nClassification Report:\n", classification_report1)
```

```
Accuracy: 0.8743932038834952
```

Confusion Matrix:

```
[[1441  0]
```

```
[ 207  0]]
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| no           | 0.87      | 1.00   | 0.93     | 1441    |
| yes          | 0.00      | 0.00   | 0.00     | 207     |
| accuracy     |           | 0.87   |          | 1648    |
| macro avg    | 0.44      | 0.50   | 0.47     | 1648    |
| weighted avg | 0.76      | 0.87   | 0.82     | 1648    |

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

### Contact Communication Type Analysis

In [8]:

```
# Assuming 'contact' is a relevant feature for analysis
```

```
X2 = sampled_data[['contact']]
```

```
y2 = sampled_data['y']
```

```
# Splitting the data into training and testing sets
```

```
from sklearn.model_selection import train_test_split
```

```
X2_train, X2_test, y2_train, y2_test = train_test_split(X2, y2, test_size=0.2, random_state=42)
```

```
# One-hot encoding for the 'contact' column
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.pipeline import Pipeline
```

```
# Define the transformer for one-hot encoding
```

```
preprocessor2 = ColumnTransformer(
```

```
    transformers=[
```

```
        ('cat', OneHotEncoder(), ['contact'])
```

```
    ],
```

```
    remainder='passthrough'
```

```
)
```

```
# Assuming Logistic Regression as the chosen classification algorithm
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
# Create a pipeline with one-hot encoding and logistic regression
```

```
model2 = Pipeline(steps=[
```

```
    ('preprocessor', preprocessor2),

    ('classifier', LogisticRegression(random_state=42))

])

# Training the model

model2.fit(X2_train, y2_train)

# Predicting responses on the test set

predictions2 = model2.predict(X2_test)

# Evaluating the model

accuracy2 = accuracy_score(y2_test, predictions2)

conf_matrix2 = confusion_matrix(y2_test, predictions2)

classification_report2 = classification_report(y2_test, predictions2)

# Displaying results

print(f'Accuracy: {accuracy2}')

print("\nConfusion Matrix:\n", conf_matrix2)
```



```
print("\nClassification Report:\n", classification_report2)
```

Accuracy: 0.8743932038834952

Confusion Matrix:

```
[[1441  0]
```

```
[ 207  0]]
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| no           | 0.87      | 1.00   | 0.93     | 1441    |
| yes          | 0.00      | 0.00   | 0.00     | 207     |
| accuracy     |           | 0.87   |          | 1648    |
| macro avg    | 0.44      | 0.50   | 0.47     | 1648    |
| weighted avg | 0.76      | 0.87   | 0.82     | 1648    |

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

### Age Groups Investigation

In [9]:

```
# Assuming 'age' and 'contact' are relevant features for analysis
```

```
X3 = sampled_data[['age', 'contact']]
```

```
y3 = sampled_data['y']
```

```
# Splitting the data into training and testing sets
```

```
from sklearn.model_selection import train_test_split
```

```
X3_train, X3_test, y3_train, y3_test = train_test_split(X3, y3, test_size=0.2, random_state=42)
```

```
# One-hot encoding for the 'contact' column
```

```
from sklearn.preprocessing import OneHotEncoder
```

```
from sklearn.compose import ColumnTransformer
```

```
from sklearn.pipeline import Pipeline
```

```
# Define the transformer for one-hot encoding
```

```
preprocessor3 = ColumnTransformer(
```

```
    transformers=[
```

```
        ('cat', OneHotEncoder(), ['contact'])
```

```
    ],
```

```
    remainder='passthrough'
```

```
)
```

```
# Assuming Logistic Regression as the chosen classification algorithm
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

*# Create a pipeline with one-hot encoding and logistic regression*

```
model3 = Pipeline(steps=[  
  
    ('preprocessor', preprocessor3),  
  
    ('classifier', LogisticRegression(random_state=42))  
  
])
```

*# Training the model*

```
model3.fit(X3_train, y3_train)
```

*# Predicting responses on the test set*

```
predictions3 = model3.predict(X3_test)
```

*# Evaluating the model*

```
accuracy3 = accuracy_score(y3_test, predictions3)
```

```
conf_matrix3 = confusion_matrix(y3_test, predictions3)
```

```
classification_report3 = classification_report(y3_test, predictions3)
```

```
# Displaying results
```

```
print(f'Accuracy: {accuracy3}')
```

```
print("\nConfusion Matrix:\n", conf_matrix3)
```

```
print("\nClassification Report:\n", classification_report3)
```

```
C:\ANA\lib\site-packages\sklearn\metrics\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
C:\ANA\lib\site-packages\sklearn\metrics\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
Accuracy: 0.8743932038834952
```

```
Confusion Matrix:
```

```
[[1441  0]
```

```
[ 207  0]]
```

```
Classification Report:
```

```
precision recall f1-score support
```

|              |      |      |      |      |
|--------------|------|------|------|------|
| no           | 0.87 | 1.00 | 0.93 | 1441 |
| yes          | 0.00 | 0.00 | 0.00 | 207  |
|              |      |      |      |      |
| accuracy     |      |      | 0.87 | 1648 |
| macro avg    | 0.44 | 0.50 | 0.47 | 1648 |
| weighted avg | 0.76 | 0.87 | 0.82 | 1648 |

C:\ANA\lib\site-packages\sklearn\metrics\\_classification.py:1248: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

In [10]:

```
#Relationship between education and responses
```

```
education_insight = sampled_data.groupby(['education', 'y']).size().unstack(fill_value=0)
```

```
print("\nInsight 1: Relationship between education and responses")
```

```
print(education_insight)
```

```
#Impact of campaign duration on responses
```

```
sampled_data['duration_category'] = pd.cut(sampled_data['duration'], bins=[0, 100, 200, 300, 400, 500, float('inf')], labels=['0-100', '101-200', '201-300', '301-400', '401-500', '501+'])
```

```
duration_insight = sampled_data.groupby(['duration_category', 'y']).size().unstack(fill_value=0)
```

```
print("\nInsight 2: Impact of campaign duration on responses")
```

```
print(duration_insight)
```

Insight 1: Relationship between education and responses

|                     | no   | yes |
|---------------------|------|-----|
| education           |      |     |
| basic.4y            | 755  | 83  |
| basic.6y            | 401  | 39  |
| basic.9y            | 1144 | 93  |
| high.school         | 1684 | 217 |
| illiterate          | 4    | 1   |
| professional.course | 943  | 124 |
| university.degree   | 2100 | 318 |
| unknown             | 272  | 60  |

Insight 2: Impact of campaign duration on responses

|   | no | yes |
|---|----|-----|
| y |    |     |



duration\_category

|         |      |     |
|---------|------|-----|
| 0-100   | 1995 | 11  |
| 101-200 | 2388 | 128 |
| 201-300 | 1277 | 155 |
| 301-400 | 696  | 116 |
| 401-500 | 366  | 88  |
| 501+    | 581  | 437 |